

Typesec

Type-Level Security for Agentic AI

Alexy and Codex

June 6, 2026

Contents

1	Preface	1
2	The Problem	2
3	The Security Model	3
4	Workspace Tour	4
5	typesec-core	4
5.1	Capabilities	5
5.2	Permissions	5
5.3	Resources	6
5.4	Policy Results	6
5.5	Minting	6
5.6	Typestate	7
5.7	Combinators	7
6	typesec-agent	7
7	RBAC	8
8	ODRL	8
9	Macros	9
10	CLI	10
11	Examples	10
11.1	RBAC Agent	10
11.2	ODRL Agent	11
11.3	Company Graph with Grust and Sail	11
11.4	Python LangChain-Style Tool Gating	11
12	Tests and Validation	12
13	Publishing the Repository	12
14	What We Improved	13
15	Design Tradeoffs	13
16	Roadmap	14
17	Conclusion	14

1 Preface

This book describes the Typesec system we designed, implemented, validated, published, and packaged today. It is both a design record and a codebase guide. The intent is not just to say what

the repository contains, but to explain why each piece exists, how the pieces fit together, and what security property the system is trying to make harder to accidentally lose.

The project began from a simple observation: agentic software makes ordinary authorization mistakes more dangerous. A human-facing web request usually has a short life. An AI agent can be long-running, can call many tools, can make many intermediate decisions, and can wander through code paths that were not anticipated when the first permission guard was written. If security depends on remembering to call `check_permission()` before every dangerous action, the system is only as strong as its most forgetful call site.

Typesec moves the central permission proof into Rust's type system. Runtime policy still matters. RBAC and ODRL policies are parsed and evaluated at runtime. But once a policy engine allows an action, the result is not just a boolean. The result is an unforgeable typed value:

```
Capability<CanWrite, Report>
```

That value becomes the proof required by sensitive APIs. Code that writes a report should accept a write capability. Code that reads sensitive data should accept a sensitive-read capability. Code that exfiltrates data should be forced to say so in its type signature. The ideal is a codebase where the compiler can show us the security boundary before the program ever runs.

The repository now lives at:

```
git@github.com:alexey/typesec.git
```

The local branch `main` tracks `origin/main`. Before publishing, the workspace was checked with:

```
cargo check --workspace
```

The Grust integration was also checked with:

```
cargo check --example company_graph_grust_sail
```

2 The Problem

Most authorization code is guard-based. A function starts with a condition:

```
if acl.check(user, "write", resource) {  
    resource.write(data);  
}
```

This pattern is familiar, easy to understand, and often good enough for small systems. It also has a structural weakness: the guard is separate from the operation. Every call path has to remember the guard. Every refactor has to preserve the guard. Every new integration has to know which guard applies.

Agentic systems stress this weakness. An agent may plan, retrieve data, call a tool, revise the plan, call another tool, ask another model, and then publish a result. Each of those steps can cross a policy boundary. Some boundaries are ordinary application boundaries, such as read versus write. Others are AI-specific boundaries, such as train, infer, or exfiltrate.

The central design question for Typesec is:

Can we make the permission proof part of the API shape, so the dangerous action is not callable unless the proof is already in scope?

Typesec does not remove runtime policy evaluation. It changes what runtime policy evaluation returns. A policy check that allows an action mints a capability. The capability is then consumed or borrowed by APIs that require that access.

The most important consequence is that permission becomes visible in function signatures:

```
fn write_report(cap: &Capability<CanWrite, Report>, report: &Report) {  
    // The body can assume the policy engine approved this subject  
    // for this permission and this resource.  
}
```

There is no matching function that takes only `&Report` and quietly writes it. If such a function appears, it is visible during review as a policy bypass.

3 The Security Model

Typesec is built around a small set of invariants.

First, capabilities are unforgeable outside `typesec-core`. The `Capability<P, R>` fields are private, and the constructor that skips policy evaluation is `pub(crate)`. External crates cannot write:

```
Capability { ... }
```

and they cannot call the unchecked constructor. They must ask a policy engine.

Second, permission types are sealed. The `Permission` trait is implemented by Typesec's own marker types, such as `CanRead`, `CanWrite`, `CanDelete`, `CanReadSensitive`, `AiCanInfer`, `AiCanTrain`, and `AiCanExfiltrate`. External crates cannot invent `CanDoAnything` and use it as a back door.

Third, resource types implement a `Resource` trait. A resource exposes a stable identifier, and that identifier is what runtime policy engines evaluate. The type parameter still matters. `Capability<CanWrite, Report>` is not the same type as `Capability<CanWrite, CustomerRecord>`, even if both are represented at runtime by strings.

Fourth, agent authentication is represented as `typestate`. The core agent starts as:

```
Agent<Unauthenticated>
```

After successful authentication it becomes:

```
Agent<Authenticated>
```

Only the authenticated state exposes capability-requesting behavior. This makes “forgot to authenticate” a type error instead of a late runtime surprise.

Fifth, every policy decision is auditable. The runtime bridge emits structured events through tracing, so production users can connect the policy layer to their logging or SIEM infrastructure.

These invariants are intentionally modest. Typesec is not pretending that the type system can know every business rule at compile time. The type system enforces that sensitive operations require a proof. Runtime policy decides whether the proof should be minted.

4 Workspace Tour

The repository is a Rust workspace with six crates:

typesec-core	traits, phantom types, Capability, PolicyEngine, tpestate
typesec-rbac	YAML RBAC parser, validator, engine, and code generator
typesec-odrl	YAML ODRL subset, constraints, prohibitions, and audit events
typesec-agent	SecureAgent wrapper for async capability requests and execute
typesec-macro	derive and policy macros for typed role declarations
typesec-cli	validate, check, generate, and run commands

The repository also includes:

```
examples/rbac_agent.rs
examples/odrl_agent.rs
examples/company_graph_grust_sail.rs
examples/langchain_company_graph.py
policies/rbac-example.yaml
policies/odrl-example.yaml
docs/company-graph-grust-sail.md
docs/improvements.md
docs/typesec-claude1.md
tests/python/test_cli_policy.py
```

The root workspace uses Rust 2024. Common dependencies are declared in [workspace.dependencies]: tokio, serde, serde_yaml, serde_json, clap, tracing, thiserror, anyhow, syn, quote, proc-macro2, and glob.

One important dependency change happened late in the day. The local example had originally depended on a path checkout of Grust:

```
grust = { path = "../../grust/crates/grust", features = ["sail"] }
```

After the Grust crates were published, the dependency was moved to the published crates:

```
[dev-dependencies]
grust-core = "0.1.0"
grust-sail = "0.1.0"
```

This was necessary because the grust crate name on crates.io is already owned by an unrelated GNOME bindings project. The property graph system is consumed through the split crates grust-core and grust-sail.

5 typesec-core

The core crate is where the security idea becomes concrete. It contains the unforgeable capability type, permission markers, resource abstraction, policy engine trait, tpestate agent, and policy combinators.

5.1 Capabilities

The central type is:

```
pub struct Capability<P: Permission, R: Resource> {
    subject: String,
    resource_id: String,
    _permission: PhantomData<fn() -> P>,
    _resource: PhantomData<fn() -> R>,
}
```

At runtime, the value stores the subject and resource identifier. The permission and resource types are represented with `PhantomData`. They cost nothing at runtime, but they force the compiler to distinguish a read capability from a write capability.

The constructor is deliberately hidden:

```
pub(crate) fn new_unchecked(
    subject: impl Into<String>,
    resource_id: impl Into<String>,
) -> Self
```

That one visibility choice carries much of the design. The capability can be minted by code inside `typesec-core`, but not by consumers. Consumers call `mint_capability`, which first asks a `PolicyEngine`.

Capabilities are intentionally not `Clone`. A privilege should not spread through the program by accident. If code needs to share a capability, it passes a reference. This is a small choice, but it aligns the ergonomics with the security model: possession of a capability is meaningful.

5.2 Permissions

Permissions are marker types. The names are runtime strings when a policy engine needs to evaluate a request, but static types when APIs require proof. The project includes ordinary permissions:

```
read
write
delete
execute
delegate
read_sensitive
write_sensitive
```

It also includes AI-oriented permissions:

```
ai:infer
ai:train
ai:exfiltrate
```

The exfiltration permission is especially important. If an agent can send data outside a trust boundary, that path should not be hidden inside an ordinary “write” operation. A function that requires `Capability<AiCanExfiltrate, CustomerData>` announces the risk in its signature.

5.3 Resources

Resources implement:

```
pub trait Resource {  
    fn resource_id(&self) -> &str;  
    fn resource_type() -> &'static str;  
}
```

The policy engine sees the identifier. Rust sees the type. This lets Typesec bridge dynamic policy files with typed application code.

For quick CLI and test paths, the core crate also provides generic resources. Examples define domain-specific resources, such as employee nodes, relationships, company graphs, and employee networks.

5.4 Policy Results

All policy engines return:

```
pub enum PolicyResult {  
    Allow,  
    Deny(String),  
    Delegate(String),  
}
```

Allow mints a capability. Deny returns an error with a reason. Delegate means this engine has no definitive answer and another engine may decide. ODRL uses delegation naturally: if an ODRL policy has no matching rule, it can defer to RBAC.

The error type for capability acquisition is:

```
pub enum CapabilityError {  
    Denied { reason: String },  
    UnhandledDelegation,  
    EngineError(String),  
}
```

One future improvement is to make delegation more first-class at the capability-minting boundary. Today, if a single engine delegates and no wrapper handles that delegation, minting fails with `UnhandledDelegation`. The composition layer solves this for deployed combinations, but the distinction is worth documenting.

5.5 Minting

The minting flow is the core runtime bridge:

```
agent.request_capability::<CanWrite, Report>(&report)  
-> engine.check(subject, "write", report.resource_id())  
-> PolicyResult::Allow  
-> Capability::new_unchecked(subject, resource_id)
```

The function that performs this is `mint_capability`. It emits an audit event for every decision and only calls the unchecked constructor after an allow.

5.6 Typestate

The typestate module defines:

```
Agent<Unauthenticated>  
Agent<Authenticated>
```

The state marker trait is sealed. External crates cannot add fake states. The only transition from unauthenticated to authenticated is through `authenticate`, which consumes the unauthenticated agent and returns an authenticated one.

The authentication implementation is intentionally small. It checks that the subject and token are present. A production system would verify a JWT, API key, or identity-provider token. The point of this crate is not to own identity. The point is to make the authenticated state visible in the type system after identity has been established.

5.7 Combinators

The combinator module lets multiple policy engines be combined. It supports four strategies:

```
PriorityOrder    first non-delegating answer wins  
AllowIfAll      all definitive engines must allow  
AllowIfAny      any allow is enough unless all deny or delegate  
DenyOverrides   any deny beats any allow
```

`DenyOverrides` is the conservative XACML-style choice. `PriorityOrder` is useful when a more specific policy should get the first chance to decide and fall back to a broader one only on delegation.

6 typesec-agent

`typesec-core` stays dependency-light. It does not need `tokio`. The `typesec-agent` crate wraps the core typestate agent in a higher-level `SecureAgent` that exposes async capability requests and execution.

The unauthenticated constructor is:

```
let agent = SecureAgent::new(Arc::new(engine));
```

Authentication returns a new type:

```
let agent = agent.authenticate(Credentials::new("agent:bot", "token"))?;
```

Only `SecureAgent<Authenticated>` has:

```
request_capability::<P, R>(&self, resource: &R)  
execute(&self, cap: &Capability<P, R>, resource: &R, action)
```

The `execute` method captures the project's main ergonomic pattern. It takes a capability reference and a resource reference, logs the execution, and then runs an async closure. The closure cannot be reached through this method without a capability.

This pattern is small enough to be adopted by application code directly. A database write wrapper can require a write capability. A vector-store retrieval wrapper can require a read capability. An outbound HTTP publisher can require an exfiltration capability.

The crate also includes `AgentBuilder`, which can wrap a single engine or build a composed engine from a primary and fallback. The default composed behavior is priority order: the primary decides unless it delegates.

7 RBAC

The `typesec-rbac` crate implements role-based access control from YAML. A policy has roles and assignments:

```
roles:
  - name: analyst
    permissions: [read, read_sensitive]
    resources: ["reports/*", "metrics/*"]
  - name: admin
    inherits: [analyst]
    permissions: [delete, delegate]
    resources: ["*"]

assignments:
  - subject: "agent:data-pipeline"
    roles: [analyst]
```

The engine compiles this into subject grants. Role inheritance is flattened during construction, not recomputed from scratch at every check. A request then checks:

```
subject -> assigned roles -> effective grants -> permission and glob match
```

The glob matching is intentionally simple. Resource patterns such as `reports/*`, `employee/**`, and `*` are enough for the first version. The code uses the `glob` crate rather than ad hoc string prefixes.

RBAC is the system's practical baseline. It answers common operational questions:

```
Can agent:data-pipeline read reports/ql?
Can agent:deploy-bot write code/main.rs?
Can agent:superuser delete anything?
```

The crate also contains code generation support. `typesec generate` can emit typed Rust structs from an RBAC policy. This is one of the paths toward the larger promise: if a role is renamed in policy, code referring to the old typed role should fail to compile.

8 ODRL

The `typesec-odrl` crate implements a focused subset of the Open Digital Rights Language. ODRL is useful for AI agents because it can express obligations, prohibitions, and contextual constraints.

A policy can say:


```

policies:
- uid: "policy:ai-agent-001"
  type: Set
  rules:
    - type: permission
      assignee: "agent:summarizer"
      action: read
      target: "asset:customer-data"
      constraints:
        - leftOperand: purpose
          operator: eq
          rightOperand: "analytics"
    - type: prohibition
      assignee: "agent:summarizer"
      action: exfiltrate
      target: "asset:customer-data"

```

This says the summarizer can read customer data for analytics, but cannot exfiltrate it. That distinction matters for agents because the same underlying data may be safe to summarize internally and unsafe to send outside the system.

The ODRL engine evaluates subject, action, target, and constraints. Constraints are evaluated against a `ConstraintContext`, which can include purpose, date-time, and custom key values. The CLI exposes `--purpose` for the common case.

Conflict resolution is conservative:

prohibition beats permission

If a matching prohibition passes its constraints, the result is a deny. If a permission matches and no prohibition overrides it, the result is allow. If no rule applies, the result is delegate.

The crate also emits structured ODRL audit events. These include the policy UID, matched rule type, subject, action, target, verdict, and constraint evaluation results. This gives operators a reason trail, not just a final boolean.

9 Macros

The `typesec-macro` crate sketches two developer-experience paths.

The first is a derive macro:

```

#[derive(TypesecRole)]
#[role(permissions = "read,write", resources = "code/*")]
pub struct Engineer;

```

The second is an inline policy macro:

```

policy! {
  role Analyst {
    can [read, read_sensitive] on ["reports/*", "metrics/*"];
  }
}

```

```
}  
}
```

Macros are not the security core. The core is capabilities and policy-engine minting. The macros are there to make typed policy declarations less tedious and to give future generated code a natural shape.

10 CLI

The CLI is the bridge for humans, scripts, and non-Rust agents. It has four commands:

validate	parse and validate a policy YAML file
check	evaluate one subject/action/resource query
generate	emit typed Rust code from an RBAC policy
run	simulate agent execution under policy

The check command is especially important because it makes Typesec usable from Python or shell without native bindings:

```
cargo run -q -p typesec-cli -- check \  
  --policy policies/rbac-example.yaml \  
  --subject agent:data-pipeline \  
  --action read \  
  --resource reports/q1
```

An allow exits successfully and prints an allow result. A deny exits with status 1. A delegate exits with status 2. That makes the command a simple policy oracle for external tools.

The format is detected from YAML shape unless `--format rbac` or `--format odrl` is passed explicitly. ODRL checks can receive a purpose:

```
cargo run -q -p typesec-cli -- check \  
  --policy policies/odrl-example.yaml \  
  --subject agent:summarizer \  
  --action read \  
  --resource customer-data \  
  --purpose analytics
```

One improvement noted in `docs/improvements.md` is to add a stable machine-readable mode, such as `typesec check --json`. Today Python can use exit codes, but parsing stdout is not ideal for long-term integrations.

11 Examples

Typesec includes examples at three layers: pure Rust RBAC, pure Rust ODRL, and company-graph integrations that show how an agent tool boundary might look in a larger system.

11.1 RBAC Agent

The RBAC example shows an authenticated agent requesting a capability and then executing a protected action. The interesting part is not the business object. The interesting part is the call shape: the protected operation is not called until after policy has minted a capability.

The example also demonstrates denial. An agent assigned to a role with read permissions cannot use a write permission unless the policy grants it.

11.2 ODRL Agent

The ODRL example demonstrates contextual policy. Purpose matters. A read for analytics may be allowed while a read for billing delegates or denies. An exfiltration action can be prohibited even if a different action on the same target is allowed.

11.3 Company Graph with Grust and Sail

The company graph example is the most complete integration. It models a company hierarchy:

```
employee:evelyn CEO
  <- REPORTS_TO employee:priya VP Engineering
    <- REPORTS_TO employee:marco Engineering Manager
      <- REPORTS_TO employee:nia Senior Software Engineer
        <- REPORTS_TO employee:omar Data Engineer
```

The policy assigns different graph-writing powers to different agents:

```
agent:executive-chief  writes executive data, company graph, and sensitive
network
agent:hr-onboarding    writes non-executive employees and reporting lines
agent:employee-nia     writes only Nia's public self-service profile
```

The Rust example uses `grust-core` to build a backend-neutral property graph and `grust-sail` to write the graph when a Sail SparkConnect server is available at `127.0.0.1:50051`. If Sail is not running, the example still demonstrates the Typesec decisions and prints the graph.

The import shape after switching to the published Grust crates is:

```
use grust_core::prelude::*;
use grust_sail::{SailConfig, SailGraphStore};
```

That change is important because it records the published-crate boundary. The example no longer depends on a sibling checkout of Grust.

11.4 Python LangChain-Style Tool Gating

The Python example is intentionally self-contained. It does not require LangChain to be installed and it does not call an LLM. Instead, it uses the shape of a LangChain tool wrapper:

```
tool call requested
-> TypesecGate.allowed(subject, action, resource)
-> cargo run -q -p typesec-cli -- check ...
-> only then mutate the graph
```

This is a realistic integration boundary. Python agents can ask the Rust CLI for a policy decision before calling a tool. The Python code treats exit code `0` as allow and any nonzero exit code as blocked. A denied tool call raises `PermissionError` before the graph mutates.

This is not as strong as Rust's type-level API, because Python cannot enforce the same phantom-type proof. But it gives Python systems a practical, testable gate today.

12 Tests and Validation

The repository has tests at several levels.

Core unit tests check that capabilities carry the right subject and resource, that read and write capability types differ, and that denied engines do not mint capabilities.

Typestate tests check that authentication transitions from unauthenticated to authenticated and that empty subjects fail.

Agent tests check the full flow:

```
construct agent
authenticate
request capability
execute with capability
```

RBAC tests cover role grants, inheritance, unknown subjects, and denials. ODRL tests cover allowed purpose, wrong purpose, prohibition, and delegation for unknown subjects.

Integration tests in `crates/typesec-agent/tests/integration.rs` exercise the agent layer across more realistic scenarios. Python smoke tests in `tests/python/test_cli_policy.py` exercise the CLI as a policy oracle.

During today's final publishing pass, the merged repository was checked with:

```
cargo check --workspace
```

When the Grust dependency was switched from a path dependency to published crates, the Grust/Sail example was checked separately:

```
cargo check --example company_graph_grust_sail
```

Both passed.

13 Publishing the Repository

The local repository initially had no remote and no commits. It was connected to GitHub as:

```
git remote add origin git@github.com:alexy/typesec.git
```

The branch was renamed from `master` to `main`, staged, and committed as:

```
dbb38aa Initial typesec workspace
```

Before pushing, a generated Python bytecode file was removed from the commit and `.gitignore` was updated with:

```
__pycache__/  
*.py[cod]
```

The remote already contained an initial commit with a `LICENSE` file:

```
2cd9182 Initial commit
```

Rather than force-push over it, the remote branch was fetched and merged:

05d8144 Merge remote-tracking branch 'origin/main'

The final push updated origin/main, and the local main branch now tracks the GitHub remote.

14 What We Improved

The improvement notes in docs/improvements.md record a useful snapshot of the engineering work.

The workspace was upgraded to Rust 2024. Compiler failures were fixed across lattice tests, proc-macro parsing, examples, and doctests. Clippy was made to pass across all targets by tightening APIs, deriving defaults, documenting public ODRL fields, and applying simplifications.

The audit integration test was stabilized with a global capture subscriber for the test binary. That avoids a racy thread-local subscriber and makes the audit story more credible.

Python smoke tests were added around typesec check, giving non-Rust agent wrappers a low-friction test path.

The company graph examples were added to show Typesec at an application boundary, not just inside small unit tests. The Rust version proves the typed capability story with Grust. The Python version proves that the CLI can gate side-effecting tools in a LangChain-like shape.

Finally, the Grust crates were switched from a local path dependency to published crates.io packages, making the example reproducible outside this machine.

15 Design Tradeoffs

Typesec deliberately separates runtime policy from compile-time proof. This is a tradeoff.

Putting all policy in types would be too rigid. Real policies come from YAML, databases, admin consoles, customer configuration, and external governance systems. They change without recompiling the application.

Leaving all policy at runtime is too easy to bypass. A forgotten guard can become a production vulnerability.

Typesec's compromise is:

```
runtime policy decides  
typed capability proves the decision happened  
protected APIs require the proof
```

Another tradeoff is that capabilities currently prove a decision at mint time. They do not yet model revocation, expiry, lease epochs, or policy-version binding. For short-lived operations, this is fine. For long-running agents, the system should grow epoch-bound or lease-bound capabilities.

The CLI is another pragmatic tradeoff. Rust code gets the strongest type-level story. Python code gets a subprocess oracle. That is weaker, but useful now. A future PyO3 crate could expose in-process policy checks, but the CLI boundary is easy to sandbox, inspect, and test.

The published Grust integration also records a naming tradeoff. The ideal facade crate name, `grust`, is unavailable on `crates.io` for this project. The example therefore imports `grust-core` and `grust-sail` directly. That is explicit and works with the published packages.

16 Roadmap

The next phase should focus on proving the central promises more directly.

First, add compile-fail tests with `trybuild`. The most important tests are:

```
unauthenticated agents cannot request capabilities
actions cannot execute without capabilities
read capabilities cannot be passed where write capabilities are required
ordinary write cannot satisfy ai:exfiltrate
```

Second, make generated policy code part of the examples. If `typesec generate` emits typed modules from an RBAC policy, downstream example code should compile against those generated types. Then a policy rename breaks code at compile time.

Third, refine deny, delegate, and constraint-failure semantics. Today ODRL uses delegation for no matching rule or failed permission constraints. That is useful for composition, but applications may want clearer distinctions in logs and errors.

Fourth, design capability expiry. A capability might carry a policy version, an epoch, or a lease deadline. Long-running agents should not keep using a proof forever after governance has changed.

Fifth, add `typesec check --json`. External agents need stable machine-readable answers:

```
{
  "verdict": "allow",
  "subject": "agent:data-pipeline",
  "action": "read",
  "resource": "reports/q1"
}
```

Sixth, deepen the Python story. The subprocess gate is enough to prove the boundary. A small Python package could wrap it cleanly. Later, a PyO3 binding could provide in-process checks for latency-sensitive integrations.

Seventh, expand the Grust example into an end-to-end backend demo that can run against a known local service profile. The current example gracefully skips `Sail` when it is not listening; a fuller demo could include setup instructions or a containerized path.

17 Conclusion

`Typesec` is a small system with a sharp idea: authorization should not only be a runtime answer. For agentic AI systems, authorization should also leave a typed trace in the program. If a function can write, read sensitive data, train a model, or exfiltrate content, that power should be visible in the function's inputs.

The repository now contains the first coherent version of that idea:

- typed capabilities
- sealed permissions
- resource abstractions
- typestate authentication
- RBAC policy evaluation
- ODRL contextual policy and prohibitions
- policy combinators
- async secure agents
- CLI policy checks
- Rust examples
- Python tool-gating example
- Grust/Sail graph integration
- tests and documentation
- published GitHub remote

The design is not finished, but it is real enough to build on. The next work is to make the compile-time guarantees more aggressively tested, make generated policy types part of ordinary workflows, and give non-Rust agents cleaner ways to use the same security boundary.

That is the arc of today's build: from an idea about impossible-to-forget authorization, to a working Rust workspace, to examples that show how agent tools can be shaped around typed proof.